

# Goals for Performance Lecture

- Understand performance, speedup, throughput, latency
- Relationship between cycle time, cycles/instruction (CPI), number of instructions (the *performance equation*)
- Amdahl's Law
- Know how to do problems in the homework!

# Definitions

Performance is in units of things-per-time

- Miles per hour, bits per second, widgets per day ...
- Bigger is better

If we are primarily concerned with response time:

- $\text{Performance}(x) = 1 / \text{ExecutionTime}(x)$

“X is n times faster than Y” means

- $n = \text{Performance}(X) / \text{Performance}(Y) = \text{Speedup}$
- If X is 1.yz times faster than Y, we can informally say that X is yz% faster than Y. Speedup is better.

# Latency vs. Throughput

## Latency (Response Time)

- How long does it take for my job to run?
- How long does it take to execute a job?
- How long must I wait for the database query?

## Throughput

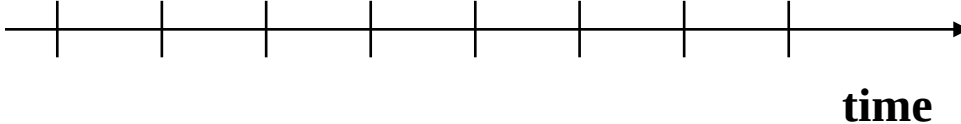
- How many jobs can the machine run at once?
- What is the average execution rate?
- How much work is getting done?

If we upgrade a machine with a new processor what do we increase?

If we add a new machine to the lab what do we increase?

# Clock Cycles

Instead of reporting execution time in seconds, we often use cycles:

$$\frac{\text{seconds}}{\text{program}} = \frac{\text{cycles}}{\text{program}} \times \frac{\text{seconds}}{\text{cycle}}$$


Clock “ticks” indicate when to start activities

Cycle time = time between ticks = seconds per cycle


Clock rate (frequency) = cycles per second (1 Hz = 1 cycle/sec)

- A 200 MHz clock has a cycle time of ...

$$\frac{1}{200 \times 10^6} \times 10^9 = 5 \text{ nanoseconds}$$

# How many cycles in a program?


Could assume that # of cycles = # of instructions



This assumption is incorrect:

- different instructions take different amounts of time on different machines (even with the same instruction set).
- Why?

# Different #s of cycles for diff'nt instrs



Multiplication takes more time than addition

Floating point operations take longer than integer ones

Accessing memory takes more time than accessing registers

Important point: changing the cycle time often changes the number of cycles required for various instructions (more later)

# Example instruction latencies

Imagine Stream Processor:

On ALU:

- Integer adds: 2 cycles
- FP adds: 4 cycles
- Logic ops (and, or, xor): 1
- Equality: 1
- < or >: 2
- Shifts: 1
- Float->int: 3
- Int->float: 4
- Select (a?b:c): 1

Other functional units:

- Integer multiply: 4
- Integer divide: 22
- Integer remainder: 23
- FP multiply: 4
- FP divide: 17
- FP sqrt: 16

# CPI

How many clock cycles, *on average*, does it take for every instruction executed?

We call this CPI (“Cycles Per Instruction”).

Its inverse ( $1/\text{CPI}$ ) is IPC (“Instructions Per Cycle”).

CISC machines: this number is

- high(er)

RISC machines: this number is

- low(er)



# CPI: Average Cycles per Instruction

$$\begin{aligned}\text{CPI} &= (\text{CPU Time} * \text{Clock Rate}) / \text{Instruction Count} \\ &= \text{Clock Cycles} / \text{Instruction Count}\end{aligned}$$

$$\text{CPI} = \sum_{i=1}^n \text{CPI}_i \times F_i \quad \text{where} \quad F_i = \frac{I_i}{\text{Instruction Count}}$$

On Imagine, integer adds are 2 cycles, FP adds are 4.  
Consider an application that has 1/3 integer adds and 2/3 FP adds.

What is its CPI?

Given a 3 GHz machine, how many instrs/sec?

# The Performance Equation

Time = Cycle Time \* CPI \* Instruction Count

- = seconds/cycle \* cycles/instr \* instrs/program
- => seconds/program

Performance = Clock Rate \* IPC \* 1/I

- = cycles/second \* instr/cycle \* program/instr
- => programs/second
  - Clock rate \* IPC = instr/second = IPS (MIPS)

“The only reliable measure of computer performance is time.”

# Remember

Performance is specific to a particular program/s

- Total execution time is a consistent summary of performance

For a given architecture performance increases come from:

- increases in clock rate (without adverse CPI affects)
- improvements in processor organization that lower CPI
- compiler enhancements that lower CPI and/or instruction count


Pitfall: expecting improvement in one aspect of a machine's performance to affect the total performance

You should not always believe everything you read! Read carefully! (see newspaper articles)

# Amdahl's Law

Speedup due to enhancement E:

$$\text{Speedup}(E) = \frac{\text{ExTime w/o } E}{\text{ExTime w/ } E} = \frac{\text{Performance w/ } E}{\text{Performance w/o } E}$$



Suppose that enhancement E accelerates a fraction F of the task by a factor S and the remainder of the task is unaffected:

$$\text{ExTime (with E)} = \frac{((1-F) + F/S) * \text{ExTime(without E)}}{1}$$

$$\text{Speedup (with E)} = \frac{1}{(1-F) + F/S}$$

Design Principle: **Make the common case fast!**

There are  
many ways to  
express  
Amdahl's Law!

# Undergrad Productivity

Average CPE student spends:

- 4 hours sleeping
- 2 hours eating
- 18 hours studying

Magic pill (purchased from dude at the UU) gives you all sleeping, eating in 1 minute!

What's the speedup on sleeping/eating?

How much more productive can you get?

# Undergrad Productivity

1

$$\text{Speedup (with E)} = \frac{1}{(1-F) + F/S}$$

F = accelerated fraction = 0.25 (6 hrs/24 hrs)

S = speedup = 6 hrs / 1 minute = 360

Overall speedup:

- $1 / [(1-0.25) + (0.25/360)]$
- $\sim 1 / (1-0.25)$
- $\sim 1.33$
- 33% more productive!