Kap. 2

Transport - Schicht

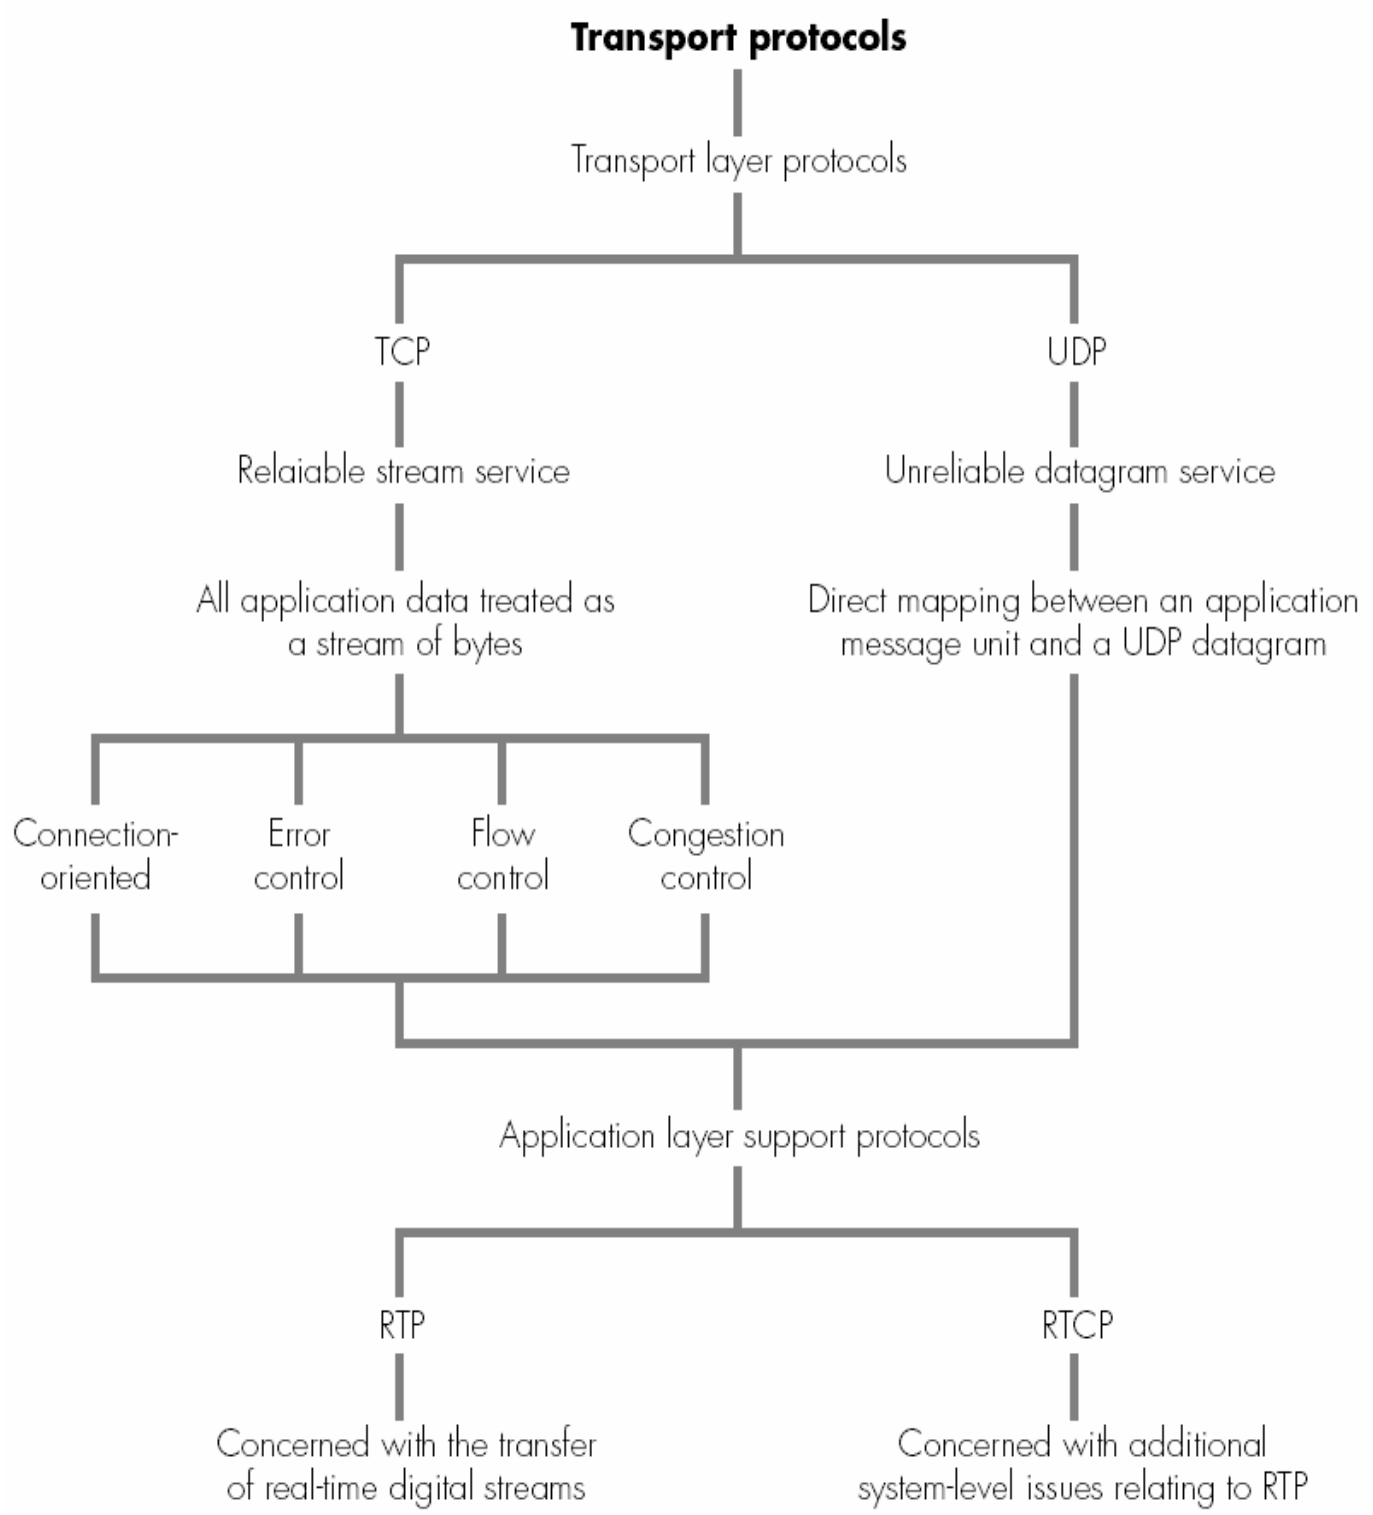
Transport protocols
Transport layer protocols
TCP
UDP
Relaiable stream service
Unreliable datagram service
All application data treated as a stream of bytes
Direct mapping between an application message unit and a UDP datagram
Connection-oriented
Error control
Flow control
Congestion control
Application layer support protocols
RTP
RTCP
Concerned with the transfer of real-time digital streams
Concerned with additional system-level issues relating to RTP

Transport-Schicht

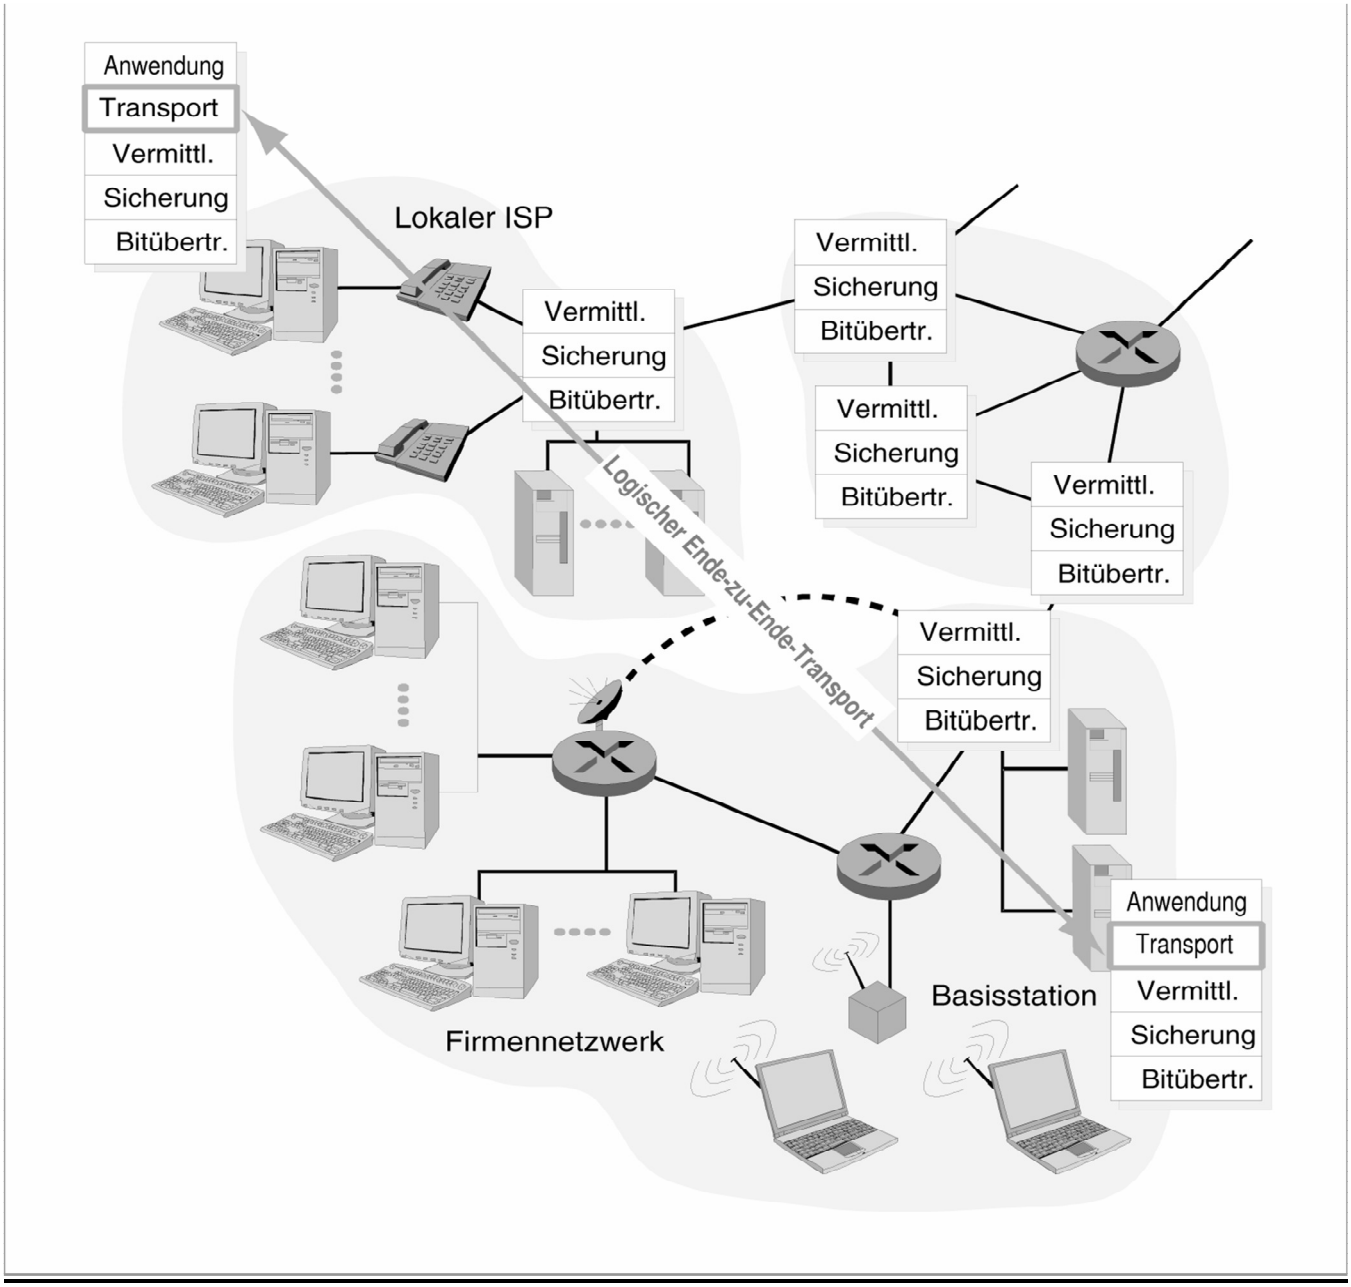

Transport-Schicht: bietet eine logische Kommunikation zw. Anwendungen

TCP: - Verbindungsorientiert mittels 3-Way-Handshake
- zuverlässiger Datentransport mittels Bestätigung
- Multiplexen & Demultiplexen von Anwendungen mittels Port-Verwaltung

UDP: - Verbindungslos
- unzuverlässige Datenübertragung
- Multiplexen + Demultiplexen von Anwendungen mittels Port-Verwaltung

- Transport protocols for the TCP/IP protocol suite:
  - *Transmission Control Protocol (TCP)*
    - Provides a connection oriented service
    - Converts the best-effort service provided by IP into a reliable service
    - Example application: file transfers
      - The transmitted information should be free of errors
      - Messages are delivered in the sequence that they are submitted
      - Example application: FTP, HTTP

  - *User Datagram Protocol (UDP)*
    - Provides a connectionless (best-effort) service
    - Example application: SNMP
      - Single request-response message exchange

- Identification of the higher layer protocol or application in the PDU header
  - IP uses the *protocol* field to identify the protocol to which the contents of the datagram relate (e.g., TCP, UDP, ICMP)
  - The transport protocol (TCP or UDP) use the *port numbers* in the PDU header to identify the application protocol to which the PDU contents relate

- Port numbers in a client-server communication
  - Client host
    - The port number of the source application protocol has only local significance
    - A new port number is allocated for each new transfer request
    - Client port numbers are called *ephemeral ports* since they are short-lived
    - Allocated in the range 1,024 through 5,000
  - Server
    - Port numbers are fixed: *well-known port numbers*
    - Allocated in the range 1 through 1,023
    - Example: file transfer (application) protocol port number is 21
    - Since a server application receives requests from multiple clients, both the source port number and the source IP address are sent to the application protocol

- Both the application and transport layers are end-to-end

- TCP provides two communicating peer application protocols (e.g., one in the client and one in the server) with a two-way, *reliable stream service*
  - Data is submitted by each local application to the TCP protocol entity as a stream of bytes
  - The stream of bytes flowing in each direction is transferred from one TCP entity to the other in a reliable way, i.e., no transmission errors, no lost or duplicate bytes

Verwendung von Ports in Client/Server Anwendungen

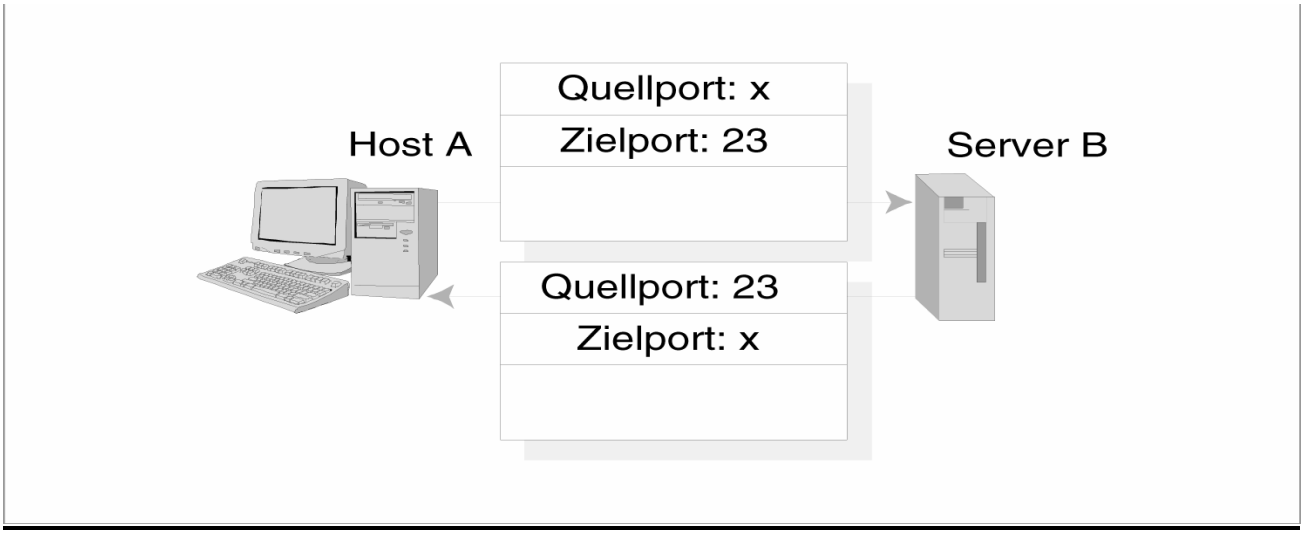

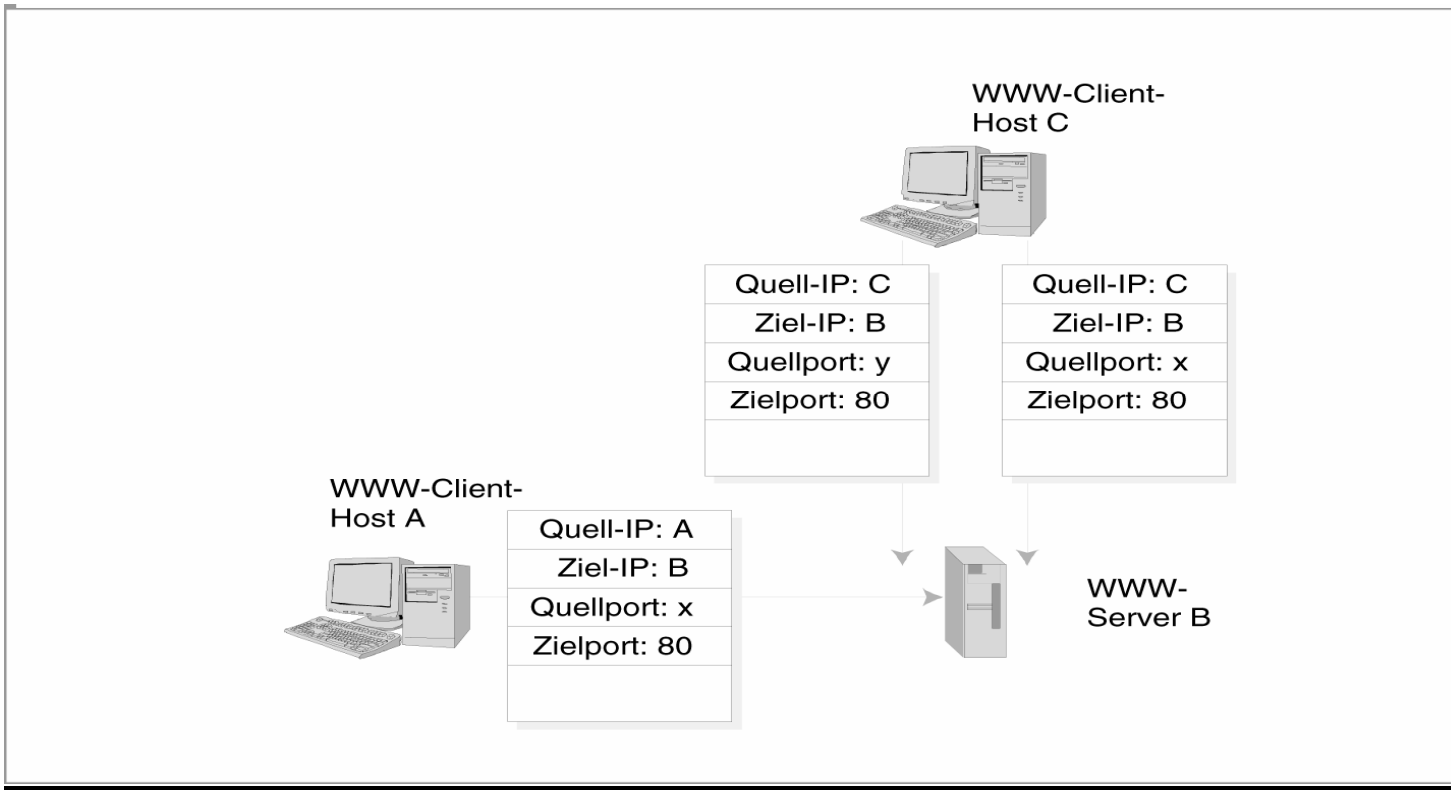

Benutzer „ A“ und „C“ : -> können die gleichen Portnr. verwenden um mit dem Server zu kommunizieren

TCP-Meldungs-Format

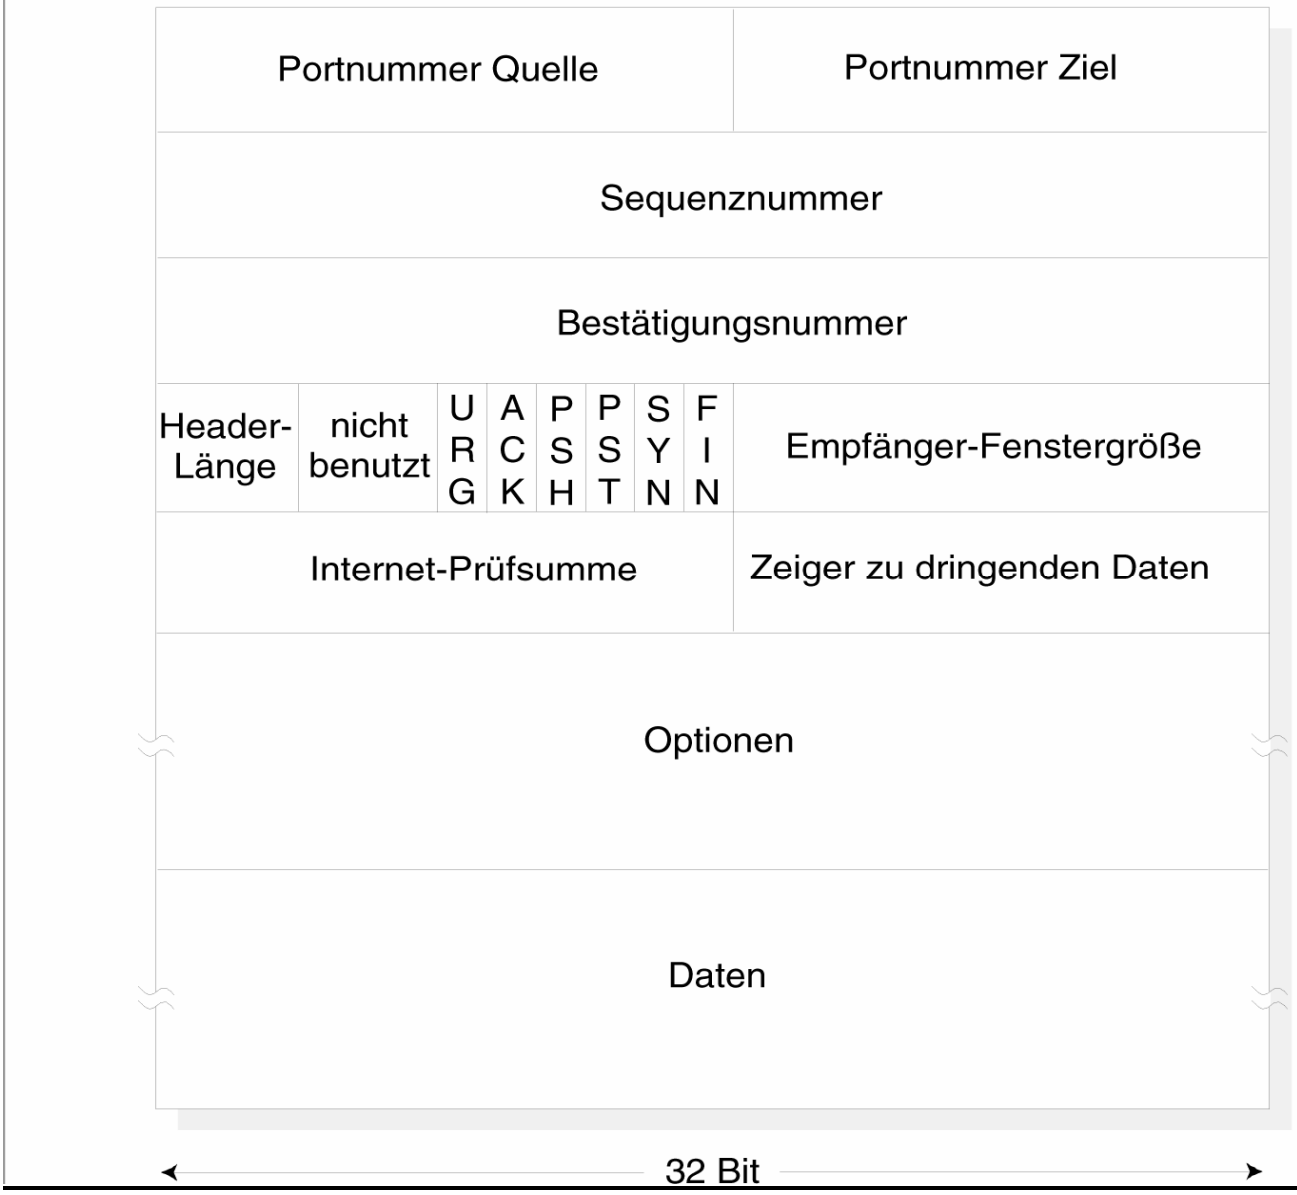

Felder:

Port Nr. Ziel/Quelle : 64.000 Ports möglich

1 – 1024 reservierte Ports

z. B.: 20, 21, 25, 80 etc., ... „well-known Ports“

Sequ. Nr. + ACK Nr: Bestätigung: z.B.:

1) Verbindungs-Auf/Abbau:
SEQnr. = ACKnr. + 1

2) Datenübertragung
SEQnr. = ACKnr. + LÄNGE in Bytes

Header-Länge: Länge des TCP-Headers; Normalerweise = 20 Bytes

FLAGS: - ***URG*** = Markiert eine Meldung als „dringend“

- *PSH* = Daten sollen sofort (auf Empfängerseite) weiter geleitet werden

- ***RST, SYN, FIN*** = Verbindungs-Auf/Abbau

- *ACK* = Bestätigung

Fenstergröße: - dient der Flusskontrolle

- # Bytes, die der Empfänger bereit ist, demnächst aufzunehmen

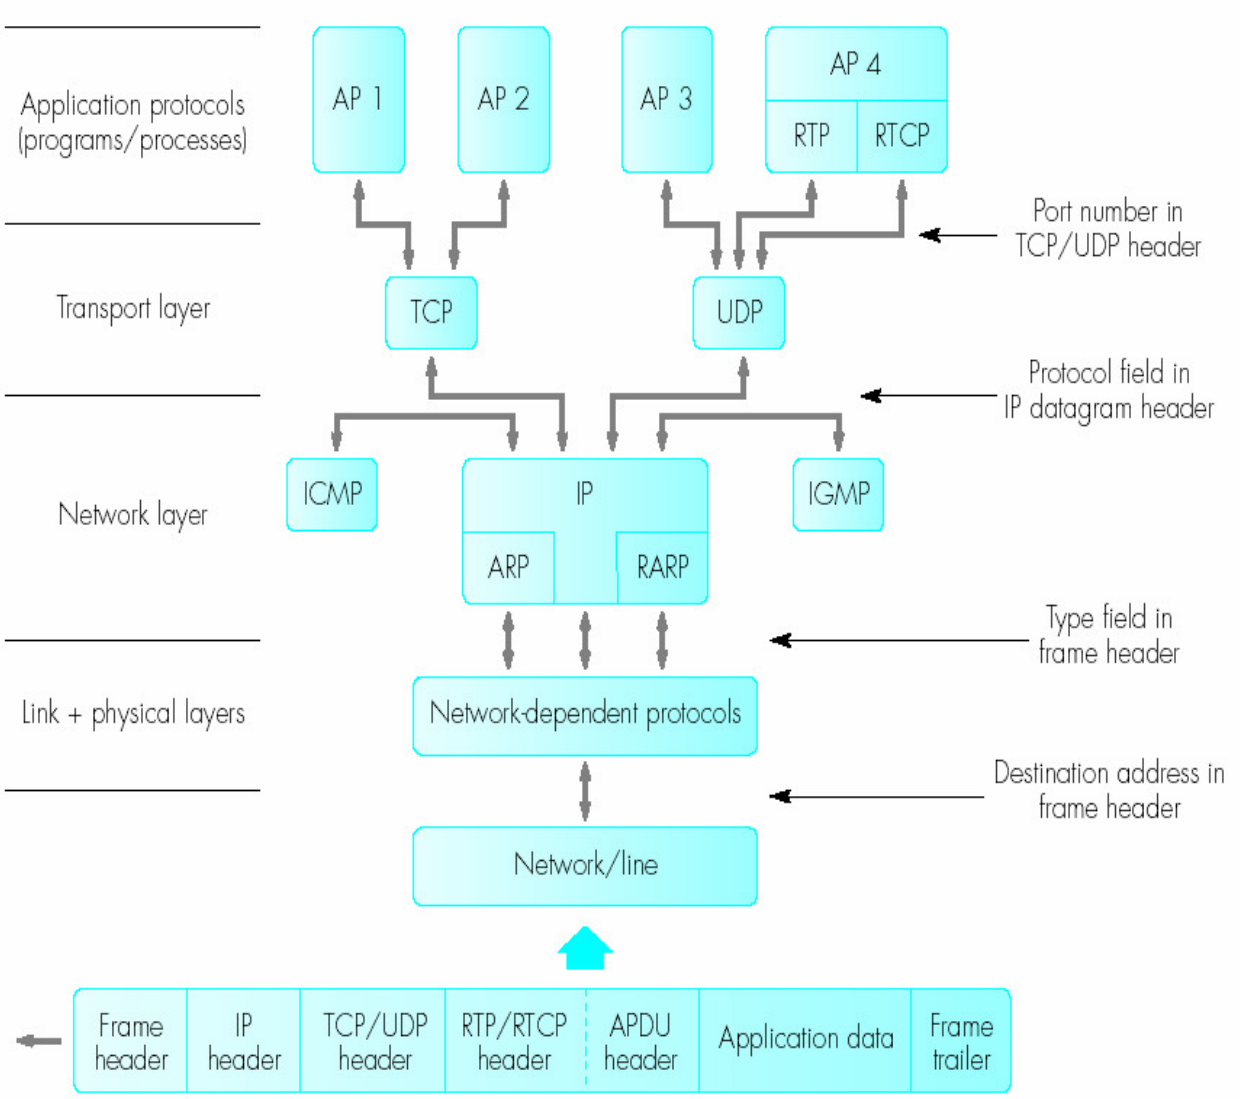
Application protocols (programs/processes)
AP 1
AP 2
AP 3
AP 4
RTP
RTCP
Port number in TCP/UDP header
Transport layer
TCP
UDP
Protocol field in IP datagram header
Network layer
ICMP
IP
IGMP
ARP
RARP
Type field in frame header
Link + physical layers
Network-dependent protocols
Destination address in frame header
Network/line
Frame header
IP header
TCP/UDP header
RTP/RTCP header
APDU header
Application data
Frame trailer

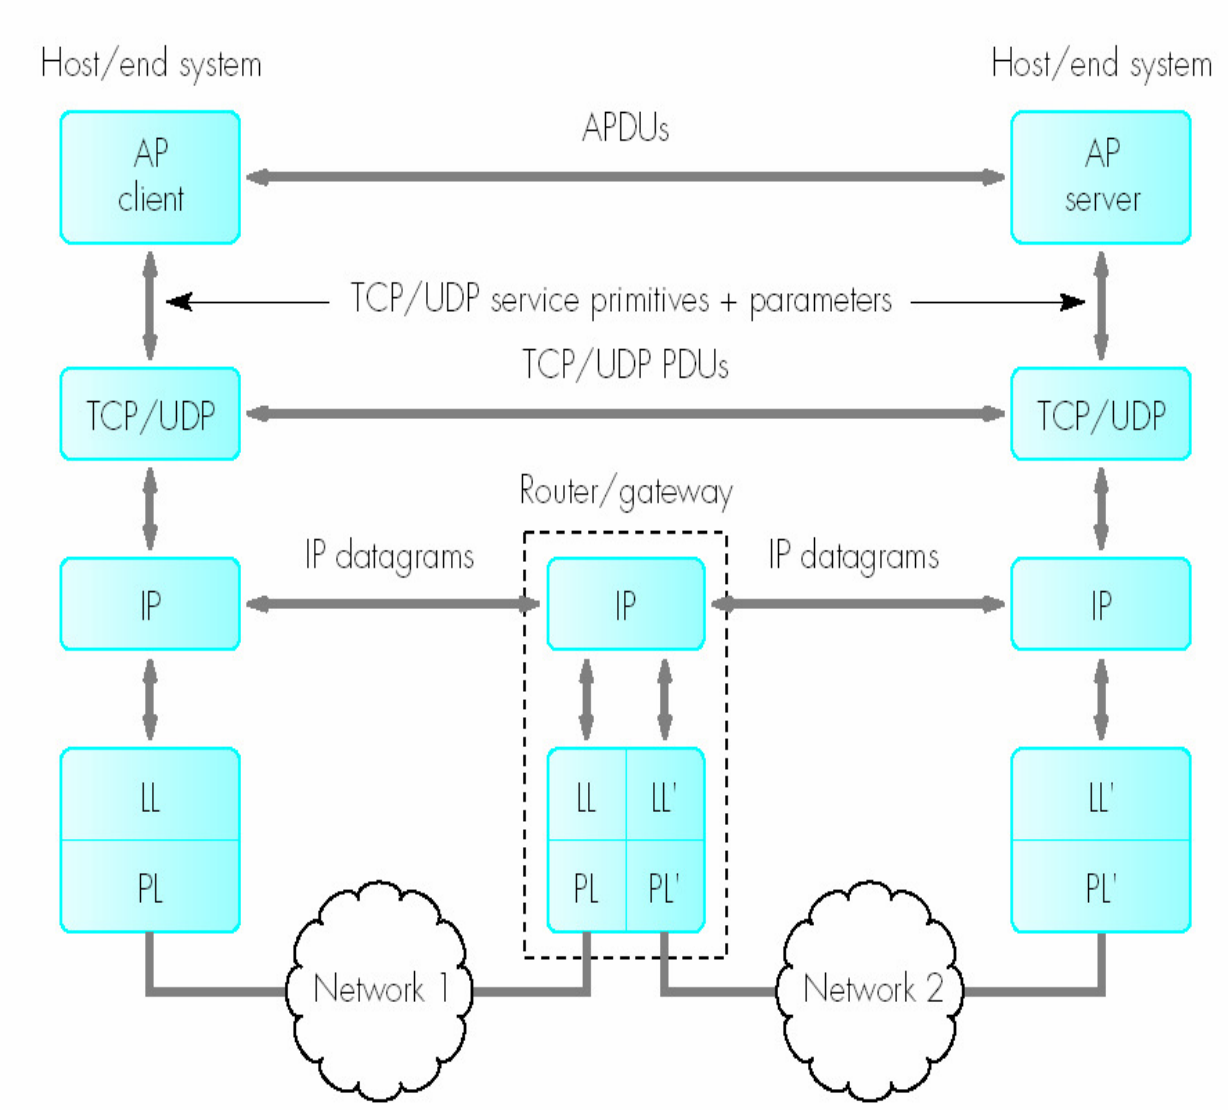
Host/end system
Host/end system
AP client
APDUs
AP server
TCP/UDP service primitives + parameters
TCP/UDP PDUs
TCP/UDP
TCP/UDP
Router/gateway
IP datagrams
IP datagrams
IP
IP
IP
LL
PL
LL
LL'
PL
PL'
LL'
PL'
Network 1
Network 2

TCP-Protokollablauf – Szenarien (1)

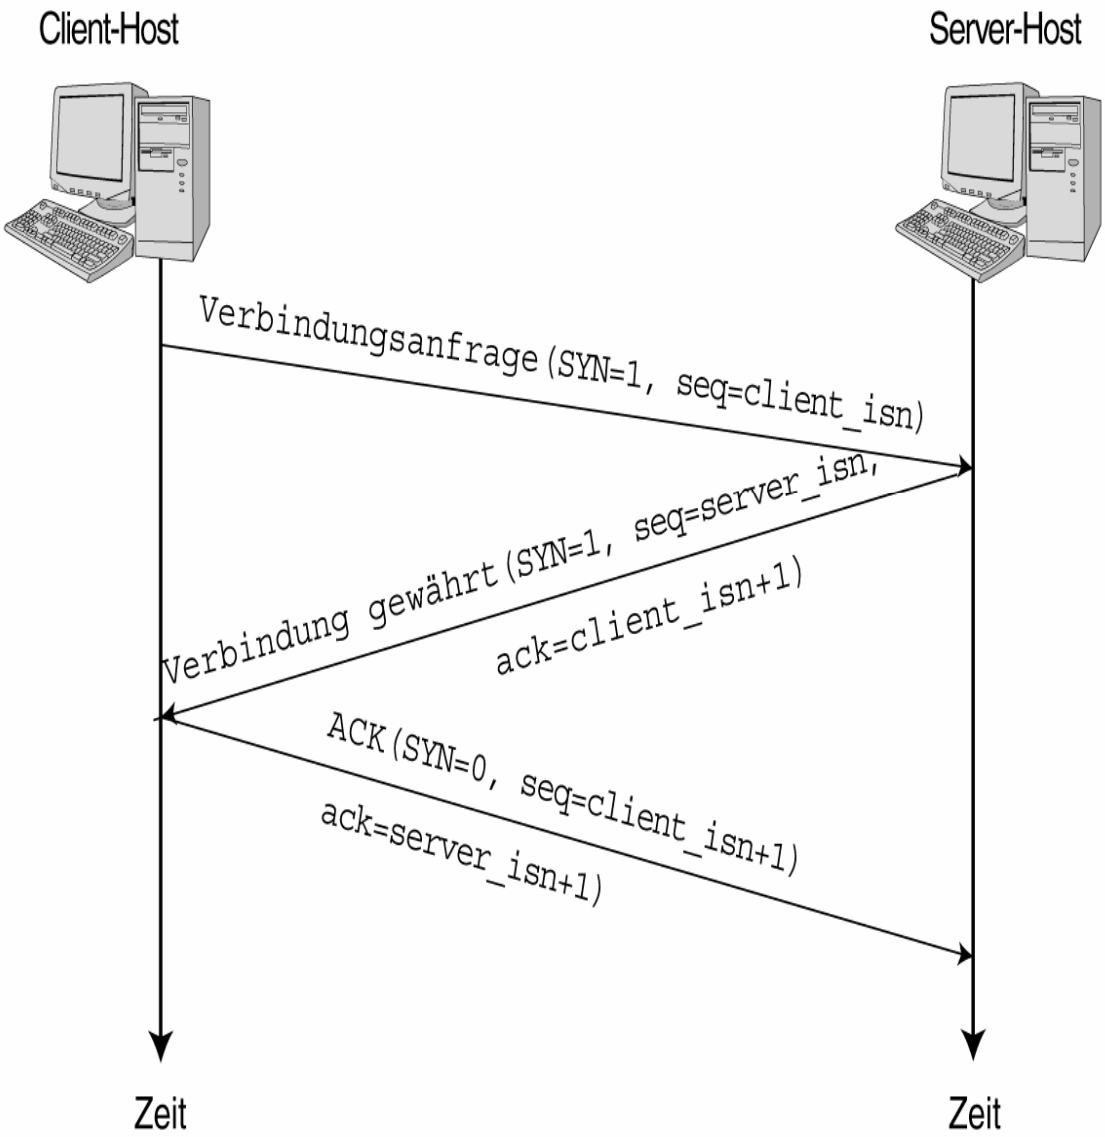

Verbindungsaufbau: 3 x Way - Handashake

TCP-Protokollablauf-Szenarien (2)

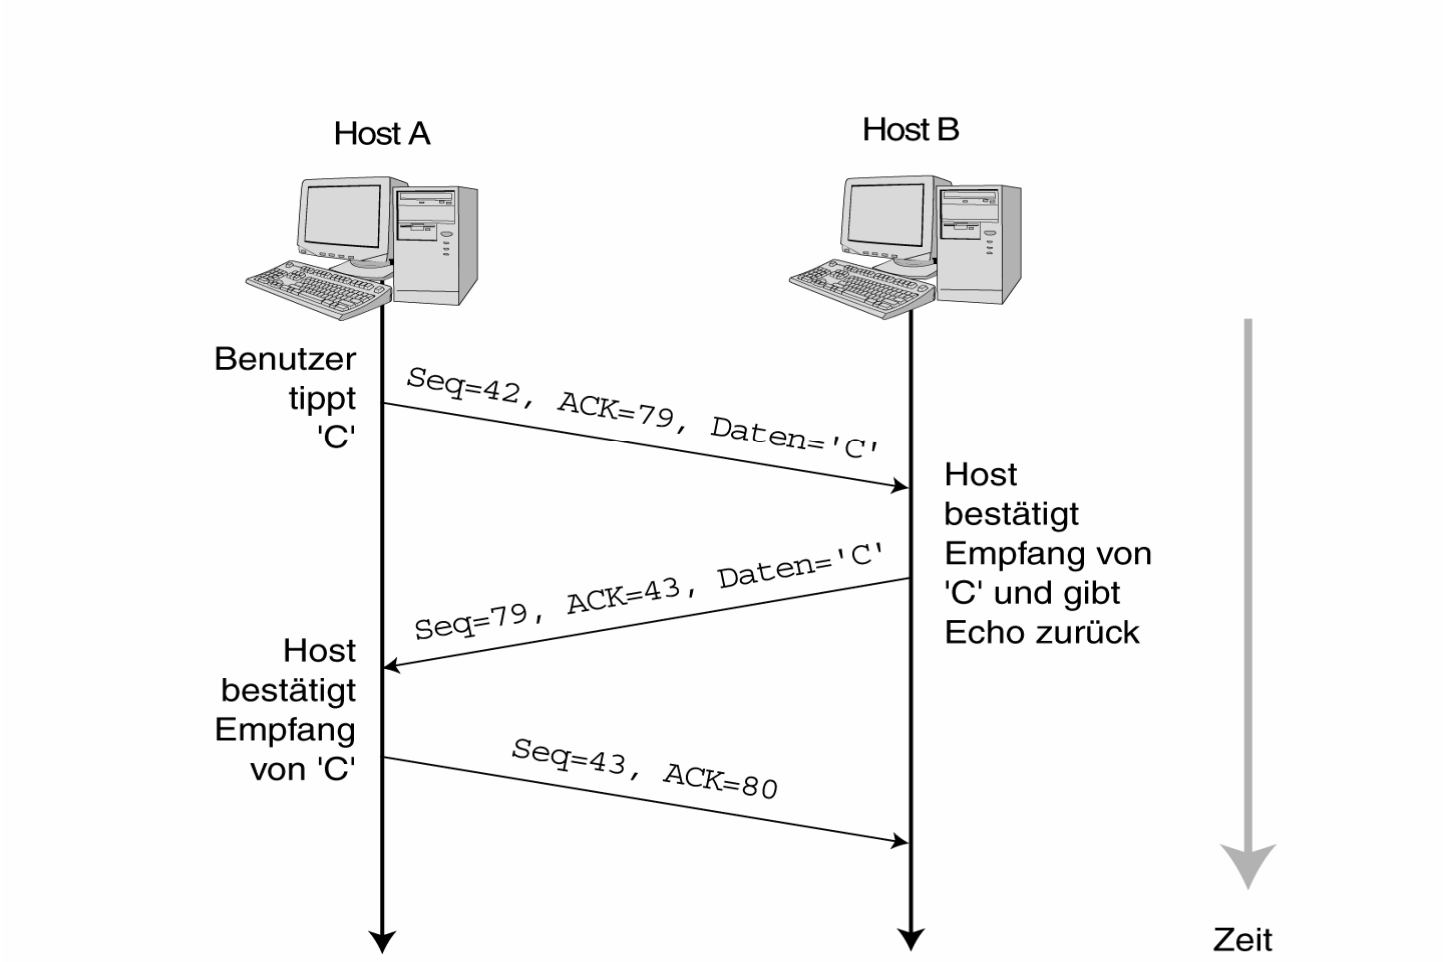

Datenübertragungs - Ablauf: SEQnr.= ACKnr. + # Bytes (Länge)

TCP-Protokollablauf-Szenarien (3)

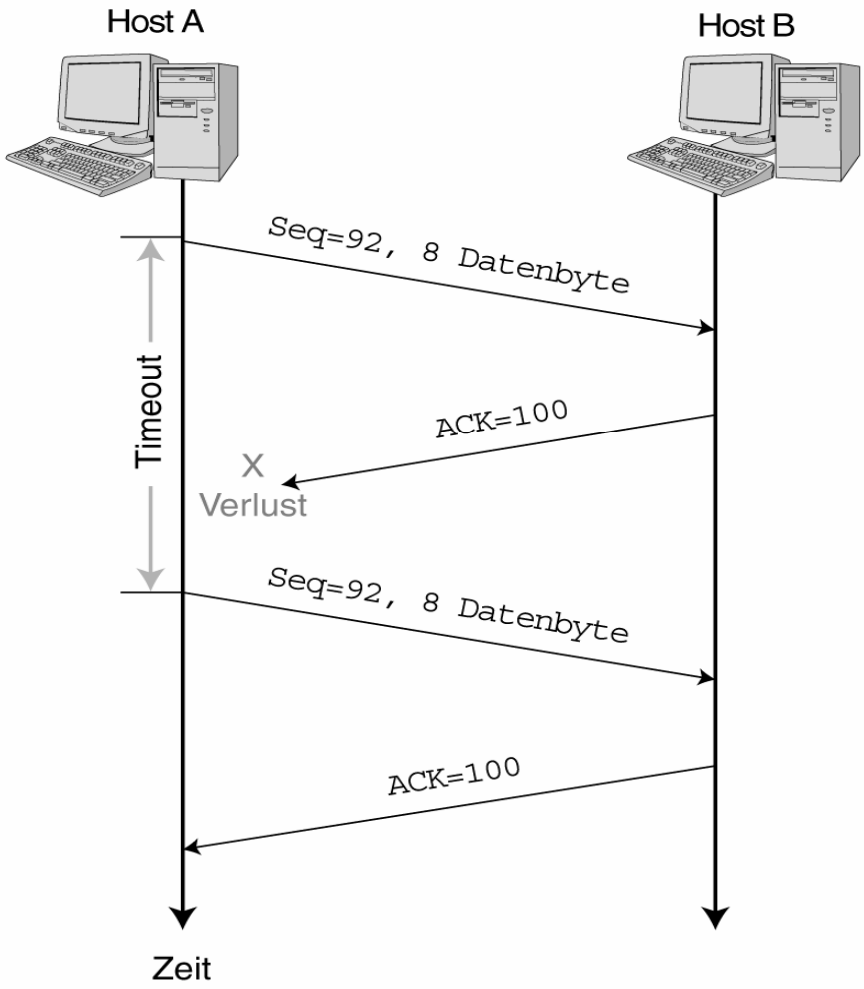

Fehlerkorrektur: Neuübertargung als Folge eines Fehlers (verlorene Bestätigung)

TCP-Protokollablauf-Szenarien (4)

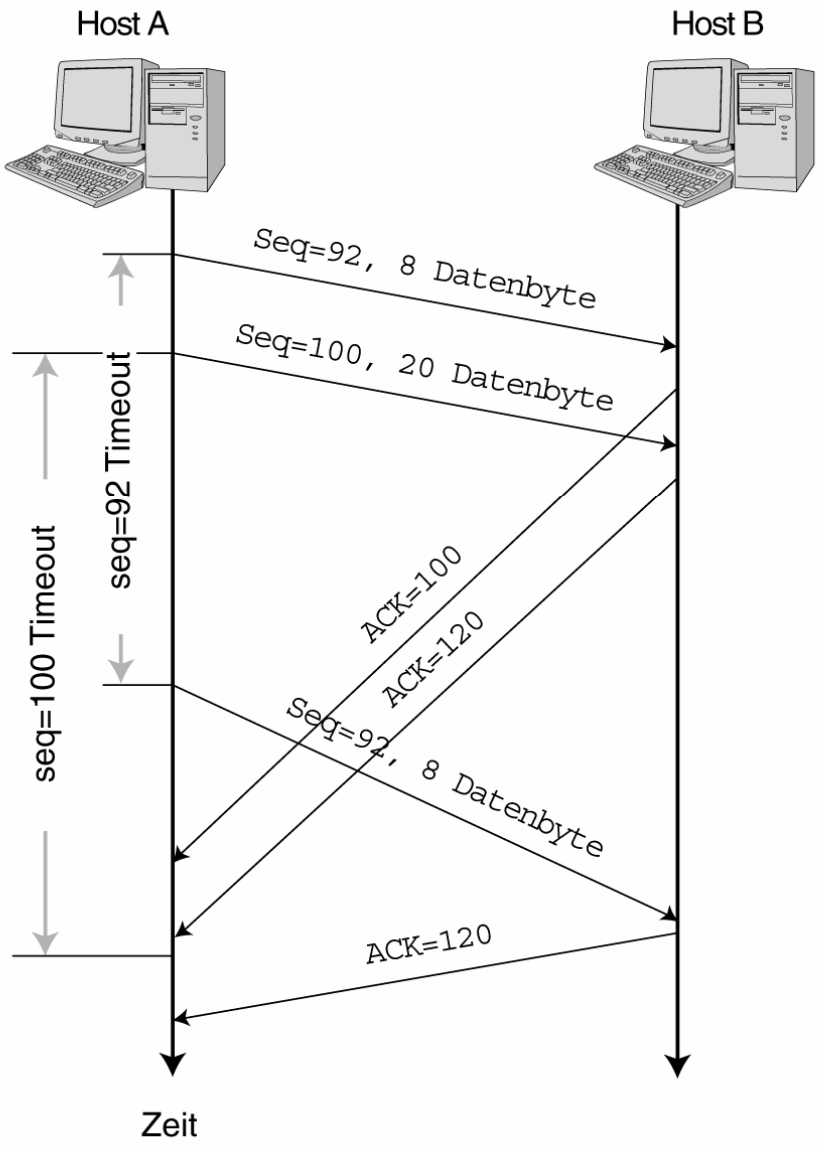

Fehlerbehandlung: Datenblock wird nicht erneut übertragen weil seine ACK vor dem Timeout ankommt

TCP-Protokollablauf-Szenarien (4)

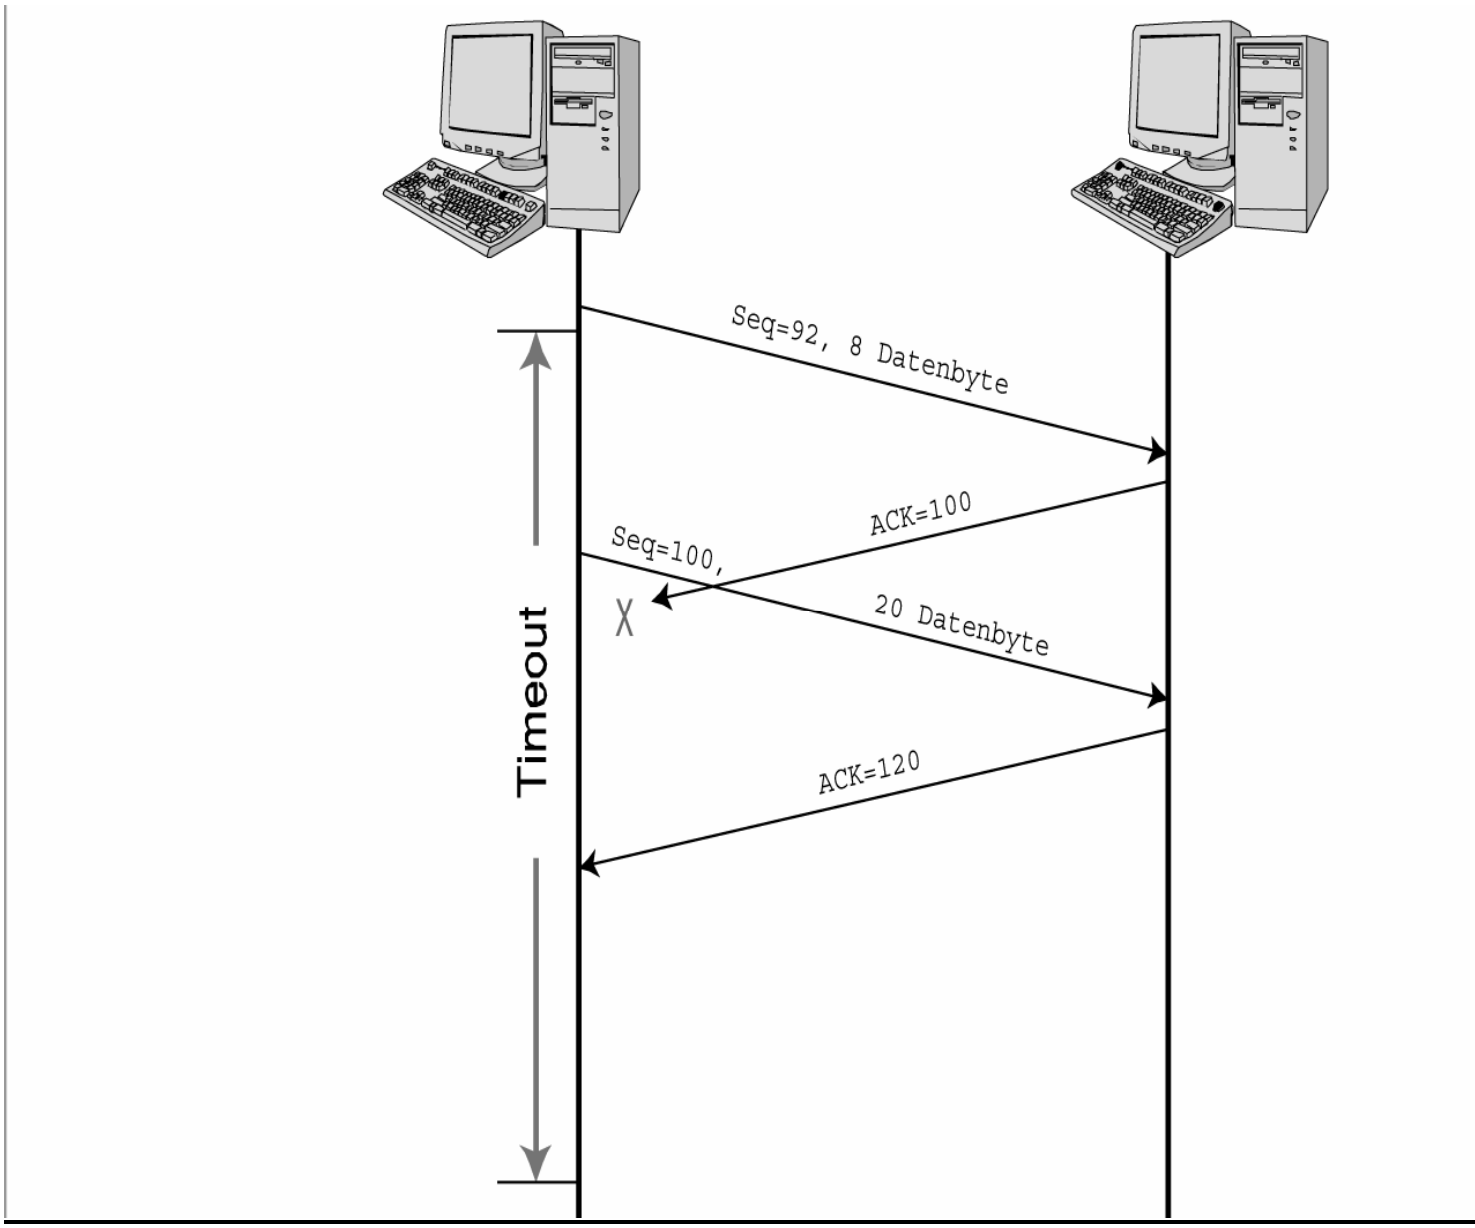

GO-BACK TO „N“ –Prinzip: Durch eine kumulative Bestätigung wird die Neuübertragung vermieden

TCP operation

- A logical connection is first established between the two TCP entities and the sequence numbers are initialized
- During the data transfer, each TCP entity divides the submitted stream of bytes into *segments*
  - A segment may contain a single byte (e.g., a character from a terminal) or many bytes (e.g., file transfer)
  - A *maximum segment size (MSS)* is agreed when the connection is set up in order to
    - Minimize transmission errors
    - Avoid fragmentation
  - Default MSS is 536 bytes
- *Flow control* is used to ensure that no data is lost when a fast server is sending data to a slower host
- *Congestion control* is used to adjust the rate of entry of segments into the network to the rate at which segments are leaving

Connection record:

- The TCP in the client may support multiple concurrent connections involving different user APs
- The TCP in the server may support multiple connections for different clients
- Data structure created for each new connection to associate received segments with the correct connection
- Contents: *Connection identifier* (pair of socket addresses), agreed *MSS*, the *initial sequence number* for each direction, the *precedence value*, and *size of window*

UDP-Protokoll: Nachrichten-Format und Funktionsweise

Port Nr. Quelle / Ziel : ähnlich wie beim TCP
Länge : Paketlänge
Prüfsumme : 1er Kompl. der Summer aller 16 Bit Wörter im Segment

Dienste von UDP: bekannt auch als Datagramm-Protokoll

- ***kein Verbindungsaufbau : keine Verzögerungen***
- ***kein Verbindungszustand***
- ***keine Bestätigung***, d. h. keine Nummerierung der Meldungen
- ***geringer Overhead*** durch Packet-Header: 8 Bytes
- Unregulierte Sende-Rate: d.h. ***keine Überlast-Kontrolle***
- Anwendungen basierend auf UDP:
  - DNS, SNMP, RIP, NF, Internet-Phone, Audio/Video-Streaming

Principles of operation

- Each message submitted by the user AP is transferred directly in a single IP datagram
- There is no connection set up, no error or flow control
- The service provided to a user AP is an extension of the service provided by IP
- The service primitives and the protocol are simpler than those of TCP
- The source UDP
  - Adds a short header to the message received from the user AP to form a *UDP datagram*
  - Forwards the UDP datagram to the IP layer for transfer over the internet using, if necessary, fragmentation
- The destination IP and UDP

- IP determines from the *protocol* field in the header that the destination protocol is UDP and passes IP datagram content to the UDP
- UDP determines the destination user AP from a field in UDP datagram header and transfers contents of datagram to that AP
- Maximum size of a UDP datagram:
  - Maximum theoretical size is 65,507 (65,535 for IP datagram – 20 bytes in IP header – 8 bytes in UDP header)
  - However, most implementations support only 8,192 bytes
  - If fragmentation is to be avoided, the size of each application PDU should be limited to the MTU for that path minus the IP and UDP headers

- - - - - - - - - - - - - - - - - Frame 5 - - - - - - - - - - - - - - - - -

SUMMARY:

	Delta T	Destination	Source	Summary
5	**0.0031**	**gigaserv.uni-..**	**[194.95.109.136)**	**DLC Ethertype=0800, size=62 bytes**
				IP D=[131.234.25.10] S=[194.95.109.136] LEN=28 ID=87
				TCP D=21 S=1036 SYN SEQ=2392861514 LEN=0 WIN=16384

DLC: ----- DLC Header -----
DLC:
DLC: Frame 5 arrived at 09:30:52.5536; frame size is 62 (003E hex) bytes.
DLC: Destination = Station 00A08E30D27F
DLC: Source = Station 00065B75C343, RFHPCI136
DLC: Ethertype = 0800 (IP)
DLC:
IP: ----- IP Header -----
IP:
IP: Version = 4, header length = 20 bytes
IP: Type of service = 00
IP: 000. .... = routine
IP: ...0 .... = normal delay
IP: .... 0... = normal throughput
IP: .... .0.. = normal reliability
IP: Total length = 48 bytes
IP: Identification = 87
IP: Flags = 4X
IP: .1.. .... = don't fragment
IP: ..0. .... = last fragment
IP: Fragment offset = 0 bytes
IP: Time to live = 128 seconds/hops
IP: Protocol = 6 (TCP)
IP: Header checksum = 2D95 (correct)
IP: Source address = [194.95.109.136]
IP: Destination address = [131.234.25.10], gigaserv.uni-paderborn.de
IP: No options
IP:
TCP: ----- TCP header -----
TCP:
TCP: Source port = 1036
TCP: Destination port = 21 (FTP)
TCP: Initial sequence number = 2392861514
TCP: Data offset = 28 bytes
TCP: Flags = 02
TCP: ..0. .... = (No urgent pointer)
TCP: ...0 .... = (No acknowledgment)
TCP: .... 0... = (No push)
TCP: .... .0.. = (No reset)
TCP: .... ..1. = SYN

TCP: ....0 = (No FIN)
TCP: Window = 16384
TCP: Checksum = B837 (correct)
TCP:
TCP: Options follow
TCP: Maximum segment size = 1460
TCP: No-op
TCP: No-op
TCP: Unknown option 4
TCP: 1 byte(s) of header padding
TCP:

Sniffer Network Analyzer data from 11-Oct-102 at 09:30:48, unsaved capture data, Page 2

- - - - - - - - - - - - - - - - Frame 6 - - - - - - - - - - - - - - - -

	Delta T	Destination	Source	Summary
6	**0.0243**	**[194.95.109.136**	**gigaserv.uni-..**	**DLC Ethertype=0800, size=62 bytes**
				IP D=[194.95.109.136] S=[131.234.25.10] LEN=28 ID=37079
				TCP D=1036 S=21 SYN ACK=2392861515 SEQ=3168076761 LEN=0 WIN=8760

DLC: ----- DLC Header -----
DLC:
DLC: Frame 6 arrived at 09:30:52.5779; frame size is 62 (003E hex) bytes.
DLC: Destination = Station 00065B75C343, RFHPCI136
DLC: Source = Station 00A08E30D27F
DLC: Ethertype = 0800 (IP)
DLC:
IP: ----- IP Header -----
IP:
IP: Version = 4, header length = 20 bytes
IP: Type of service = 00
IP: 000. .... = routine
IP: ...0 .... = normal delay
IP: .... 0... = normal throughput
IP: .... .0.. = normal reliability
IP: Total length = 48 bytes
IP: Identification = 37079
IP: Flags = 4X
IP: .1.. .... = don't fragment
IP: ..0. .... = last fragment
IP: Fragment offset = 0 bytes
IP: Time to live = 246 seconds/hops
IP: Protocol = 6 (TCP)
IP: Header checksum = 2714 (correct)
IP: Source address = [131.234.25.10], gigaserv.uni-paderborn.de
IP: Destination address = [194.95.109.136]

IP: No options
IP:
TCP: ----- TCP header -----
TCP:
TCP: Source port   = 21 (FTP)
TCP: Destination port   = 1036
TCP: Initial sequence number = 3168076761
TCP: Acknowledgment number   = 2392861515
TCP: Data offset   = 28 bytes
TCP: Flags   = 12
TCP: ..0. .... = (No urgent pointer)
TCP: ...1 .... = Acknowledgment
TCP: .... 0... = (No push)
TCP: .... .0.. = (No reset)
TCP: .... ..1. = SYN
TCP: ....0 = (No FIN)
TCP: Window   = 8760
TCP: Checksum   = 1540 (correct)
TCP:
TCP: Options follow
TCP: No-op
TCP: No-op
TCP: Unknown option 4
TCP: 5 byte(s) of header padding
TCP:

Sniffer Network Analyzer data from 11-Oct-102 at 09:30:48, unsaved capture data, Page 3

- - - - - - - - - - - - - - - - - Frame 7 - - - - - - - - - - - - - - - - -

	Delta T	Destination	Source	Summary
7	**0.0001**	**gigaserv.uni-..**	**[194.95.109.136**	**DLC Ethertype=0800, size=60 bytes**
				IP D=[131.234.25.10] S=[194.95.109.136]
				LEN=20 ID=88
				TCP D=21 S=1036 ACK=3168076762
				WIN=17520

DLC: ----- DLC Header -----
DLC:
DLC: Frame 7 arrived at 09:30:52.5780; frame size is 60 (003C hex) bytes.
DLC: Destination = Station 00A08E30D27F
DLC: Source   = Station 00065B75C343, RFHPCI136
DLC: Ethertype = 0800 (IP)
DLC:
IP: ----- IP Header -----
IP:
IP: Version = 4, header length = 20 bytes
IP: Type of service = 00
IP: 000. .... = routine
IP: ...0 .... = normal delay
IP: .... 0... = normal throughput

IP: .... .0.. = normal reliability
IP: Total length = 40 bytes
IP: Identification = 88
IP: Flags = 4X
IP: .1.. .... = don't fragment
IP: ..0. .... = last fragment
IP: Fragment offset = 0 bytes
IP: Time to live = 128 seconds/hops
IP: Protocol = 6 (TCP)
IP: Header checksum = 2D9C (correct)
IP: Source address = [194.95.109.136]
IP: Destination address = [131.234.25.10], gigaserv.uni-paderborn.de
IP: No options
IP:
TCP: ----- TCP header -----
TCP:
TCP: Source port = 1036
TCP: Destination port = 21 (FTP)
TCP: Sequence number = 2392861515
TCP: Acknowledgment number = 3168076762
TCP: Data offset = 20 bytes
TCP: Flags = 10
TCP: ..0. .... = (No urgent pointer)
TCP: ...1 .... = Acknowledgment
TCP: .... 0... = (No push)
TCP: .... .0.. = (No reset)
TCP: .... ..0. = (No SYN)
TCP: ....0 = (No FIN)
TCP: Window = 17520
TCP: Checksum = 1FCC (correct)
TCP: No TCP options
TCP:

- - - - - - - - - - - - - - - - Frame 48 - - - - - - - - - - - - - - - -

	Delta T	Destination	Source	Summary
48	**0.0012**	**gigaserv.uni-..**	**[194.95.109.1..**	**DLC Ethertype=0800, size=60 bytes**
				IP D=[131.234.25.10] S=[194.95.109.136] LEN=20 ID=109
				TCP D=21 S=1036 FIN ACK=3168079573
				SEQ=2392861579 LEN=0 WIN=16822

DLC: ----- DLC Header -----
DLC:
DLC: Frame 48 arrived at 09:31:02.9334; frame size is 60 (003C hex) bytes.
DLC: Destination = Station 00A08E30D27F
DLC: Source = Station 00065B75C343, RFHPCI136
DLC: Ethertype = 0800 (IP)
DLC:
IP: ----- IP Header -----
IP:
IP: Version = 4, header length = 20 bytes
IP: Type of service = 00
IP: 000. .... = routine
IP: ...0 .... = normal delay
IP: .... 0... = normal throughput
IP: .... .0.. = normal reliability
IP: Total length = 40 bytes
IP: Identification = 109
IP: Flags = 4X
IP: .1.. .... = don't fragment
IP: ..0. .... = last fragment
IP: Fragment offset = 0 bytes
IP: Time to live = 128 seconds/hops
IP: Protocol = 6 (TCP)
IP: Header checksum = 2D87 (correct)
IP: Source address = [194.95.109.136]
IP: Destination address = [131.234.25.10], gigaserv.uni-paderborn.de
IP: No options
IP:
TCP: ----- TCP header -----
TCP:
TCP: Source port = 1036
TCP: Destination port = 21 (FTP)
TCP: Sequence number = 2392861579
TCP: Acknowledgment number = 3168079573
TCP: Data offset = 20 bytes

TCP: Flags = 11
TCP: ..0. .... = (No urgent pointer)
TCP: ...1 .... = Acknowledgment
TCP: .... 0... = (No push)
TCP: .... .0.. = (No reset)
TCP: .... ..0. = (No SYN)
TCP: ....1 = FIN
TCP: Window = 16822
TCP: Checksum = 174A (correct)
TCP: No TCP options
TCP:

Sniffer Network Analyzer data from 11-Oct-102 at 09:30:48, unsaved capture data, Page 2

- - - - - - - - - - - - - - - - - Frame 49 - - - - - - - - - - - - - - - - -

	Delta T	Destination	Source	Summary
49	**0.0066**	**[194.95.109.1..**	**gigaserv.uni-..**	**DLC Ethertype=0800, size=60 bytes**
				IP D=[194.95.109.136] S=[131.234.25.10] LEN=20 ID=37099
				TCP D=1036 S=21 FIN ACK=2392861579 SEQ=3168079573 LEN=0 WIN=8760

DLC: ----- DLC Header -----
DLC:
DLC: Frame 49 arrived at 09:31:02.9401; frame size is 60 (003C hex) bytes.
DLC: Destination = Station 00065B75C343, RFHPCI136
DLC: Source = Station 00A08E30D27F
DLC: Ethertype = 0800 (IP)
DLC:
IP: ----- IP Header -----
IP:
IP: Version = 4, header length = 20 bytes
IP: Type of service = 10
IP: 000. .... = routine
IP: ...1 .... = low delay
IP: .... 0... = normal throughput
IP: .... .0.. = normal reliability
IP: Total length = 40 bytes
IP: Identification = 37099
IP: Flags = 4X
IP: .1.. .... = don't fragment
IP: ..0. .... = last fragment
IP: Fragment offset = 0 bytes
IP: Time to live = 246 seconds/hops
IP: Protocol = 6 (TCP)
IP: Header checksum = 26F8 (correct)
IP: Source address = [131.234.25.10], gigaserv.uni-paderborn.de
IP: Destination address = [194.95.109.136]
IP: No options
IP:

TCP: ----- TCP header -----

TCP:
TCP: Source port = 21 (FTP)
TCP: Destination port = 1036
TCP: Sequence number = 3168079573
TCP: Acknowledgment number = 2392861579
TCP: Data offset = 20 bytes
TCP: Flags = 11
TCP: ..0. .... = (No urgent pointer)
TCP: ...1 .... = Acknowledgment
TCP: .... 0... = (No push)
TCP: .... .0.. = (No reset)
TCP: .... ..0. = (No SYN)
TCP: ....1 = FIN
TCP: Window = 8760
TCP: Checksum = 36C8 (correct)
TCP: No TCP options
TCP: